

**Fill out the Brown
Computer Science Survey
you got in your email!**

Only takes 5
min!

All multiple
choice!

o / The
o Percentage
o Project 2020
percentageproject.com

*If you didn't receive the survey, email
litofish@cs.brown.edu*

FEB
13

Challenging Silicon Valley's Infinite Loop of Irresponsibility

Free

Register



Brown alumna and New York Times reporter Natasha Singer will talk about a new movement -- "responsible computing"

About this Event

Date And Time

Thu, February 13, 2020
4:00 PM – 5:30 PM EST
[Add to Calendar](#)

Sets, Dictionaries & Hash Tables

CS16: Introduction to Data Structures & Algorithms
Spring 2020

Q: how would you build a (basic) search engine?

What's so Hard about Search Engines?

"The **Google** Search **index** contains
hundreds of billions of webpages and
is well over 100,000,000 gigabytes in
size."

How Google Search Works | Crawling &
Indexing

<https://www.google.com/search/crawl...>

The screenshot shows a Google search results page. At the top, the Google logo is on the left, and the search query "Brown University" is in the search bar on the right, accompanied by a magnifying glass icon. Below the search bar, a navigation bar includes tabs for "All", "News", "Images", "Maps", "Videos", "More", "Settings", and "Tools". The "All" tab is underlined. A status message "About 3,730,000,000 results (1.05 seconds)" is displayed. The first search result is for "Brown University", with the URL <https://www.brown.edu/>. A red box highlights the search time "(1.05 seconds)". Below the result, a snippet of text reads: "Brown University, founded in 1764, is a member of the Ivy League and recognized for the quality of its teaching, research, and unique curriculum. Providence ...".

Search Through Each Page?

- ▶ Assume Google indexes 200 billion pages
- ▶ If we scan 1 page in 1 microsecond
 - ▶ each search would take 55 hours
- ▶ How can we improve search time?




Outline

- ▶ Sets
- ▶ Dictionaries
- ▶ Hash Tables
- ▶ Ex: Search engine



Dictionary

- ▶ Collection of key/value pairs
 - ▶ distinct and unordered keys
- ▶ Supports value lookup by key
- ▶ Also known as a *map*
 - ▶ “maps” keys to values
- ▶ examples
 - ▶ name → address
 - ▶ word → definition



Dictionary ADT

- ▶ **add(key, value):**
 - ▶ adds key/value pair to dict.
- ▶ **object get(key):**
 - ▶ returns value mapped to key
- ▶ **remove(key):**
 - ▶ removes key/value pair
- ▶ **int size():**
 - ▶ returns number key/value pairs
- ▶ **boolean isEmpty():**
 - ▶ returns TRUE if dict. is empty; FALSE otherwise

Q: how can we implement a dictionary?

Array-based Dictionary

- ▶ Can we use an expandable array A ?
- ▶ **add(k, v):**
 - ▶ store (k, v) in first empty cell of A
 - ▶ takes $O(1)$ if you keep track of first empty cell
- ▶ **get(k):**
 - ▶ scan A to find value with key $\text{key}=k$
 - ▶ takes $O(n)$
- ▶ **remove(k):**
 - ▶ scan A to find pair with $\text{key}=k$ & remove
 - ▶ takes $O(n)$

Is **$O(n)$** good enough? What if our dictionary stores 200B key/value pairs?

Q: can we do better?

Yes! with a Hash Table

- ▶ Hash tables are composed of
 - ▶ an array A
 - ▶ and a “hash” function $h: X \rightarrow Y$

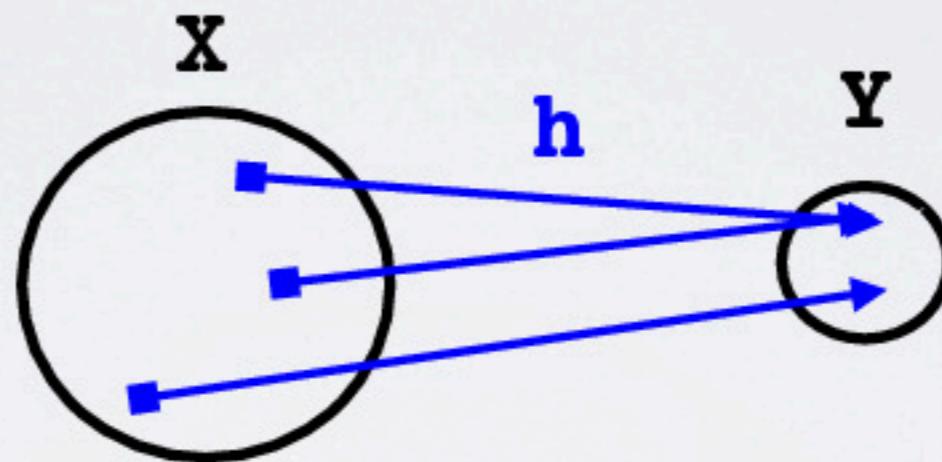


Dictionary vs. Hash Table

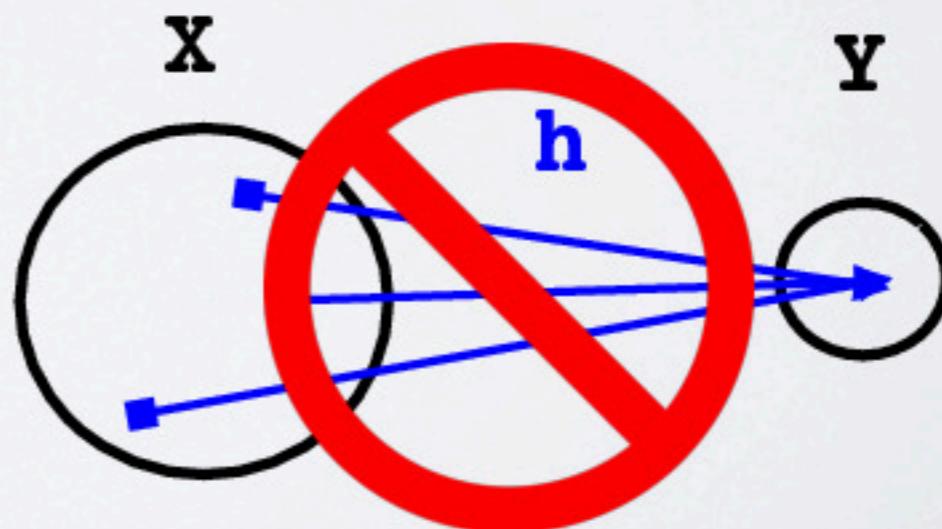
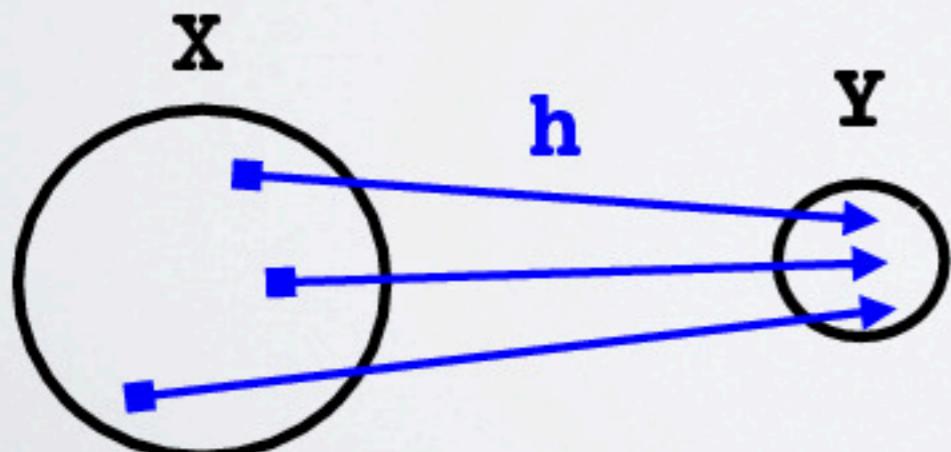
- ▶ A dictionary (or map) is an *abstract data type*
 - ▶ can be implemented using many different *data structures*
- ▶ A hash table is a *dictionary data structure*
 - ▶ one specific way to implement a dictionary

Yes! with a Hash Table

- ▶ A hash function is function $\mathbf{h}: \mathbf{X} \rightarrow \mathbf{Y}$ that
 - ▶ shrinks: maps elements from a large input space to a smaller output space



- ▶ well spread: \mathbf{h} spreads elements of \mathbf{X} over \mathbf{Y}



Building a Dictionary w/ a Hash Table



- ▶ Choose a hash function $h : X \rightarrow Y$ with
 - ▶ X = “universe of keys” and Y = “indices of array”
 - ▶ **add(k, v)**
 - ▶ set $A[h(k)] = v$ which is $O(1)$
 - ▶ **get(k)**
 - ▶ return $v = A[h(k)]$ which is $O(1)$
 - ▶ **remove(k)**
 - ▶ delete $A[h(k)]$ which is $O(1)$

HashTable — Add

keys: banner IDs

values: names

00943855

Kaila Jeter

00745911

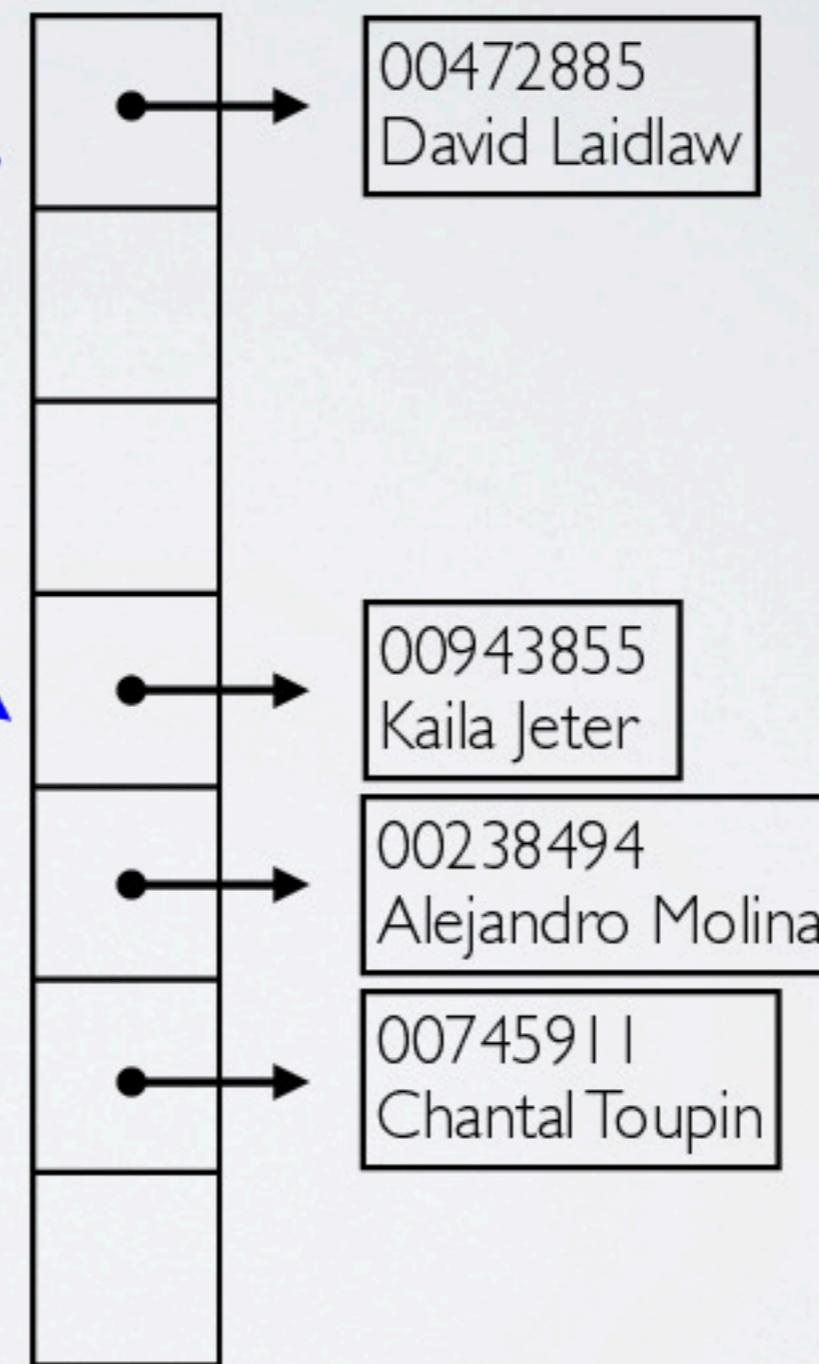
Chantal Toupin

00238494

Alejandro
Molina

00472885

David Laidlaw



Building a Dictionary w/ a Hash Table

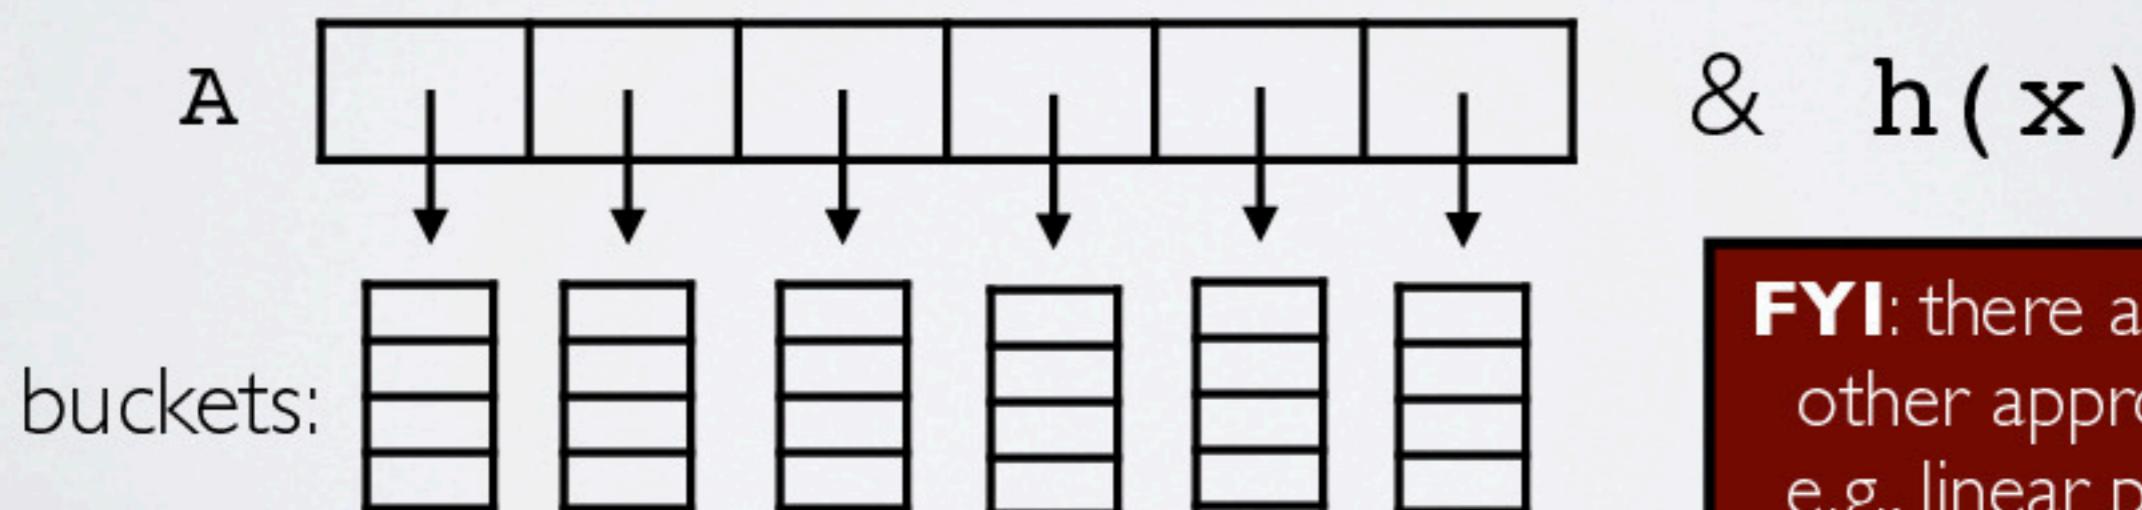


- ▶ **Q:** What is the problem with this?
 - ▶ Remember that $|Y| < |X|$
 - ▶ (here $|X|$ denotes size of X)
 - ▶ ...so some keys in X will be hashed to the same location!
 - ▶ this is called the *pigeonhole principle*
 - ▶ there just isn't enough room in Y to fit all of X
 - ▶ ...therefore some values in array will be overwritten
 - ▶ this is called a *collision*

Overcoming Collisions



- ▶ Hash Table with *Chaining*
 - ▶ store *multiple* values at each array location
 - ▶ each array cell stores a “bucket” of pairs
 - ▶ can implement bucket as a list or expandable array or ...



FYI: there are many other approaches e.g., linear probing, quadratic probing, cuckoo hashing,...

HashTable

```
table: array  
h: hash function
```

```
function add(k, v):  
    index = h(k)  
    table[index].append(k, v)
```

$O(1)$ if computing
hash function
is $O(1)$

```
function get(k):  
    index = h(k)  
    for (key, val) in table[index]:  
        if key = k:  
            return val  
    error("key not found")
```

runtime
depends on
bucket size

Hash Table

- ▶ Let's do another example but with Chaining!
- ▶ We'll use the following hash function
 - ▶ $h(\text{banner_id}) = \text{banner_id} \% 7$

Hash Table — Add

keys: banner IDs

values: names

00943855

Kaila Jeter

00745911

Chantal Toupin

00238494

Alejandro
Molina

00472885

David Laidlaw

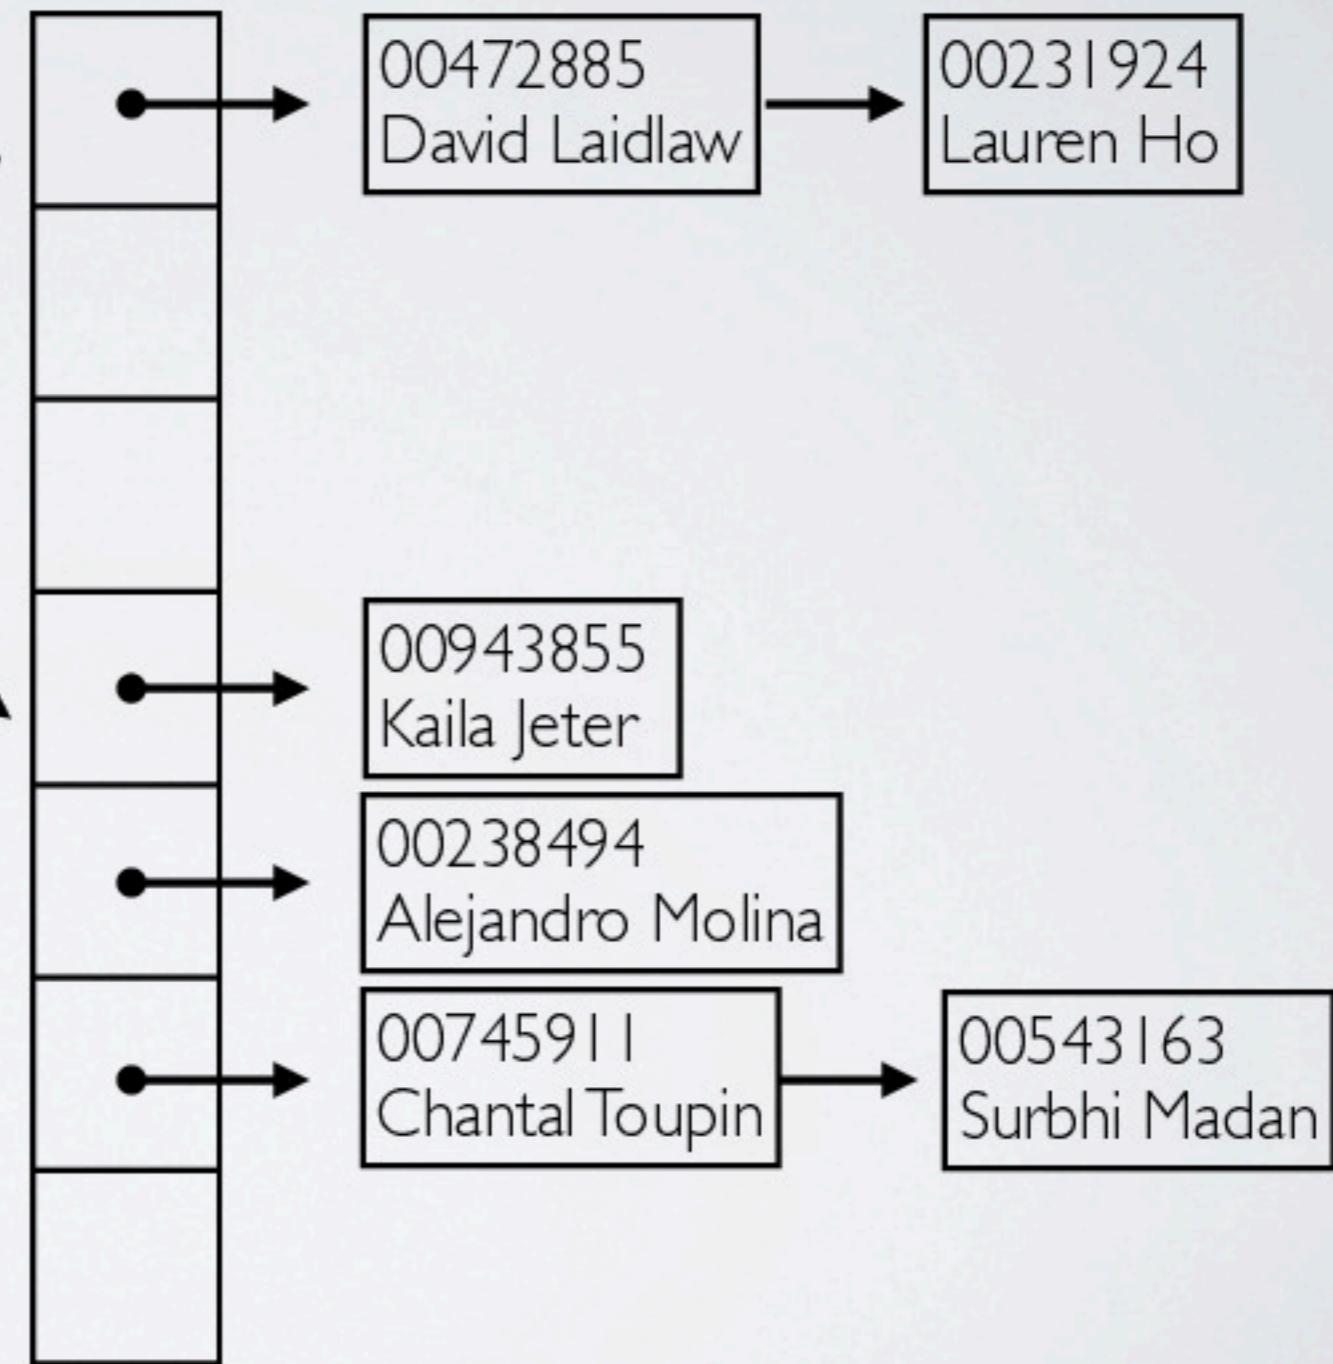
00231924

Lauren Ho

00543163

Surbhi Madan

$$h(\text{key}) = \text{key} \% 7$$



HashTable — Get

keys: banner IDs

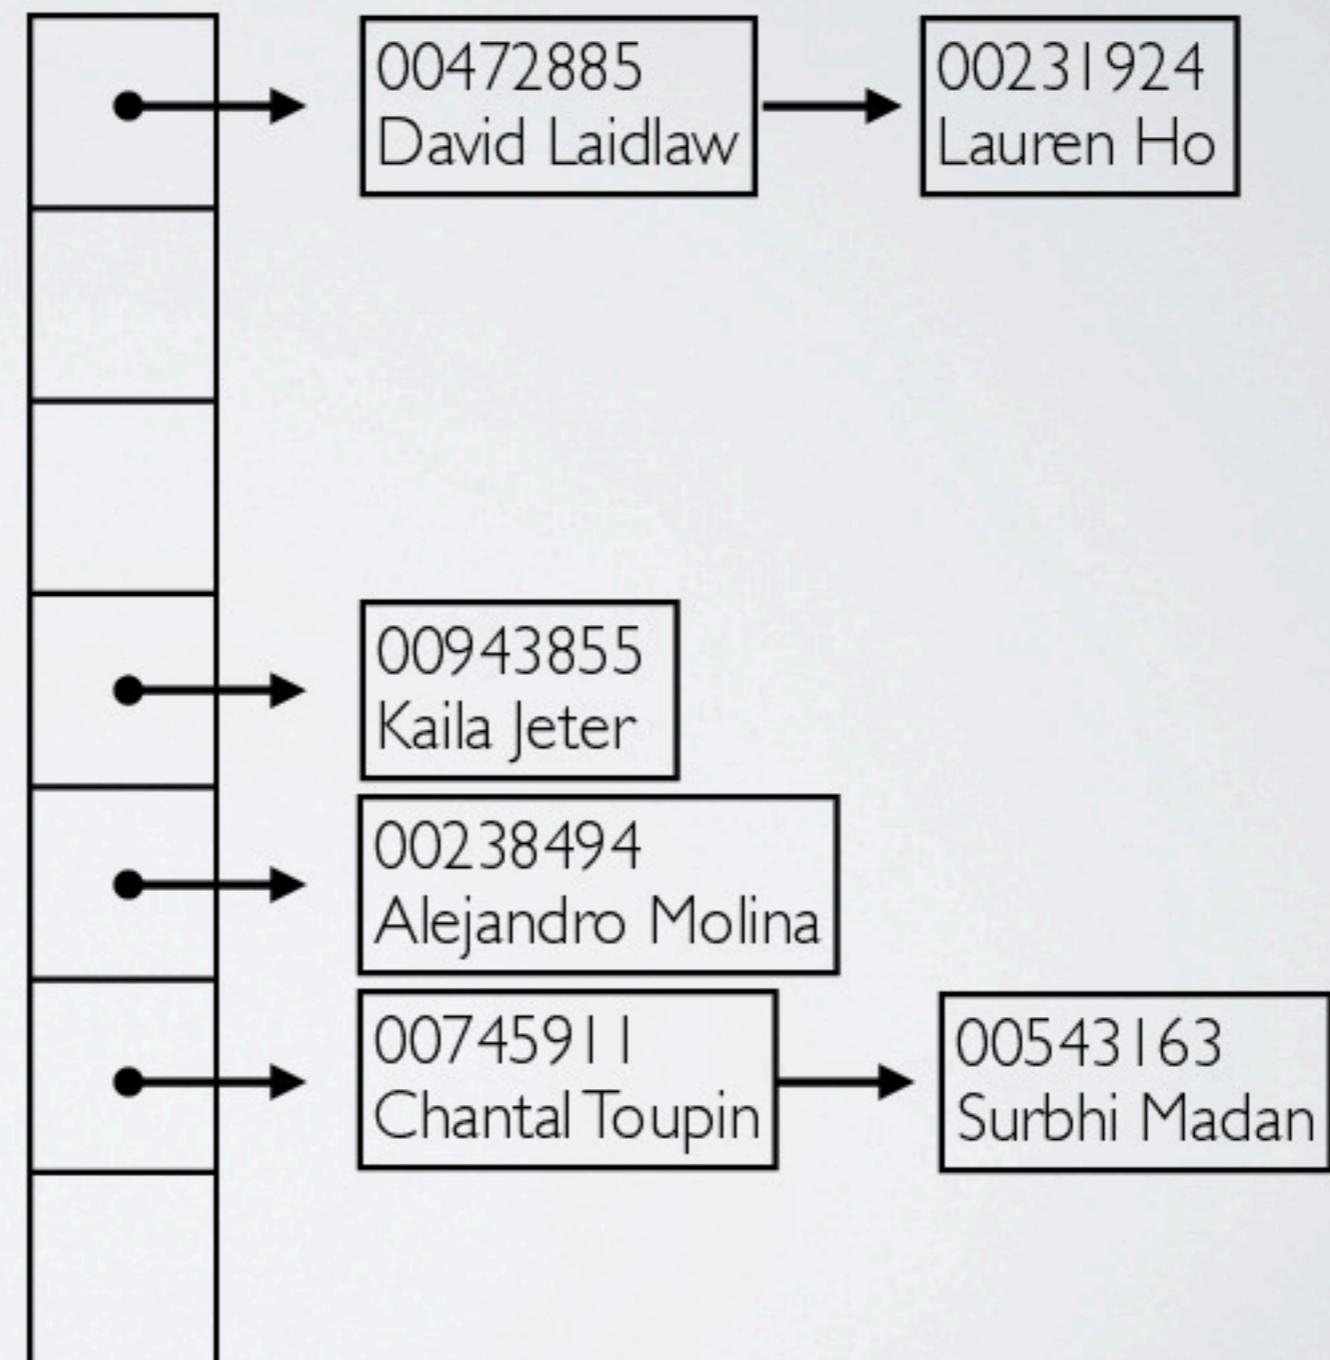
values: names

$$h(\text{key}) = \text{key} \% 7$$

00543163

**What is the
worst-case
run time of Get?**

**Array of buckets w/
key/value pairs**



HashTable with Chaining

- ▶ What is the worst-case runtime of Get?
 - ▶ \approx size of largest bucket
- ▶ What is the size of largest bucket?
 - ▶ assume we have **n** students and a table of size **m**
 - ▶ if **h** “spreads” keys roughly evenly then
 - ▶ each bucket has size $\approx n/m$
 - ▶ ex: if **n=150** and **m=7** each buckets has size $\approx 150/7 = 21$
- ▶ But what is the size of the largest bucket *asymptotically*?
 - ▶ assume **m** is a constant (i.e., it does not grow as a function of **n**)
 - ▶ each bucket has size $\approx n/m = n/c = O(n)$ 

Q: Can we do better than $O(n)$?

Beating $O(n)$ — Idea #1



- ▶ **Idea:** use large table
- ▶ Banner IDs have 8 digits so max ID is 99,999,999
- ▶ Use table of size $m=100,000,000$
 - ▶ w/ hash function $h(key) = key$
- ▶ Are there any collisions in this case?
 - ▶ no collisions because every pair gets its own cell
 - ▶ What is run time of Get?
 - ▶ $O(1)$ since we don't need to scan buckets
- ▶ What is the problem with this approach?
 - ▶ what if we only store 150 students? we're wasting 99,999,850 cells

Beating $O(n)$ — Idea #2

- ▶ **Idea:** use a table of size equal to the number of students + “good” hash function
 - ▶ set the table size to $m=n$
 - ▶ use a hash function h that spreads keys well
- ▶ No wasted space since $n = m$
 - ▶ in other words, “table size” = “number of students”
- ▶ If h spreads keys roughly evenly then each bucket has size
 - ▶ $\approx n/m = n/n = 1 = O(1)$
- ▶ What hash function should we use?
 - ▶ Suppose $n = 150$ (i.e., we want to insert 150 students)
 - ▶ should we use the hash function $h(\text{key}) = \text{key} \% 150$?

Banner ID Hashing

Form groups of 10

5 min

Activity #1

Banner ID Hashing

5 min

Activity #1

Banner ID Hashing

Activity #1

4 min

Banner ID Hashing

3 min

Activity #1

Banner ID Hashing

2 min

Activity #1

Banner ID Hashing

1 min **Activity #1**

Banner ID Hashing

Omin **Activity #1**

Beating $O(n)$ — Idea #2



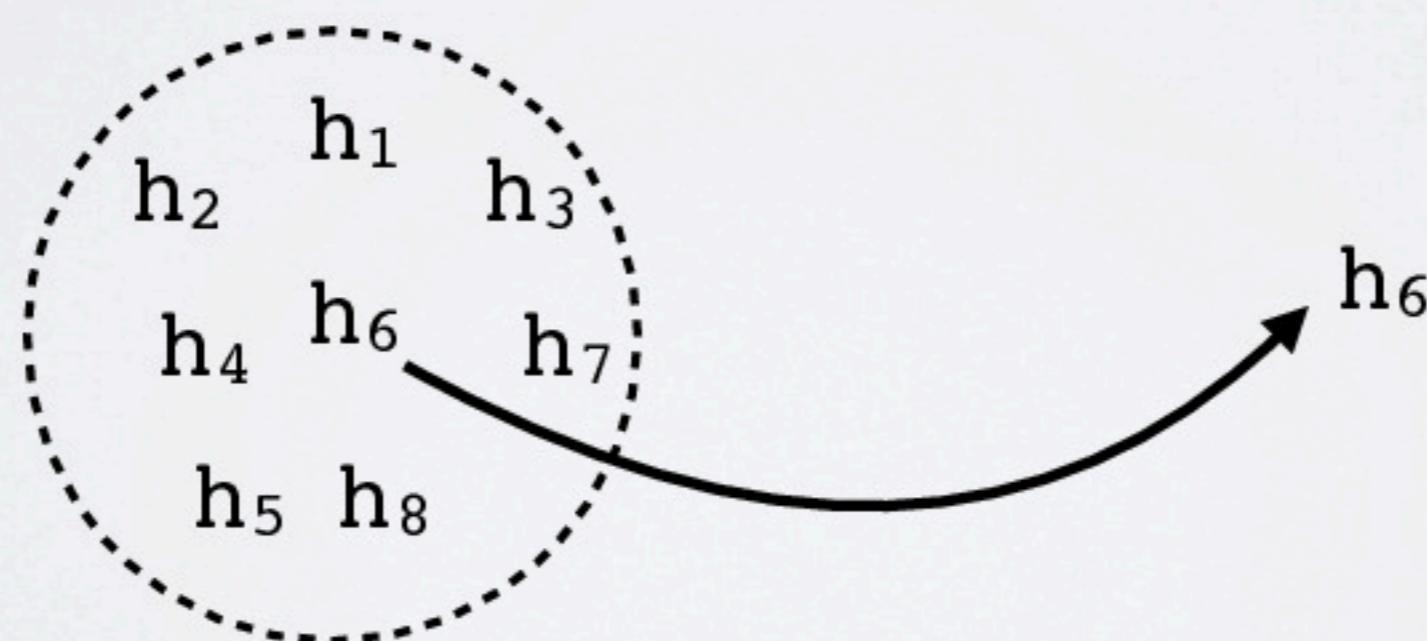
- ▶ Idea #2 relied on an assumption:
 - ▶ ***if h spreads*** keys roughly evenly then each bucket has size
 - ▶ $\approx n/m = n/n = 1 = O(1)$
- ▶ Will $h(ID) = ID \% 11$ spread banner IDs evenly?
 - ▶ it depends on the banner IDs...
 - ▶ if banner IDs are chosen randomly then Yes
 - ▶ But what if next year all banner IDs are multiples of 11?
 - ▶ Then *all* banner IDs will map to 0!
 - ▶ So there will be one bucket with all IDs
 - ▶ so worst-case runtime of Get will be $O(n)$



**Since keys are not necessarily random, we
make the hash function random**

Universal Hash Functions

- ▶ Special “families” of hash functions
 - ▶ **UHF** = $\{h_1, h_2, \dots, h_q\}$
 - ▶ designed so that if we pick a function from the family at random and use it on a set of keys, then it is very *likely* that the function will “spread” the keys (roughly) evenly



Universal Classes of Hash Functions

J. LAWRENCE CARTER AND MARK N. WEGMAN

IBM Thomas J. Watson Research Center, Yorktown Heights, New York 10598

Received August 8, 1977; revised August 10, 1978

This paper gives an *input independent* average linear time algorithm for storage and retrieval on keys. The algorithm makes a random choice of hash function from a suitable class of hash functions. Given any sequence of inputs the expected time (averaging over all functions in the class) to store and retrieve elements is linear in the length of the sequence. The number of references to the data base required by the algorithm for any input is extremely close to the theoretical minimum for any possible hash function with randomly distributed inputs. We present three suitable classes of hash functions which also can be evaluated rapidly. The ability to analyze the cost of storage and retrieval without worrying about the distribution of the input allows as corollaries improvements on the bounds of several algorithms.

Example of Universal Hash Functions

- ▶ Setup to store **n** key/value pairs
 - ▶ choose prime **p** larger than **n**
 - ▶ choose **4** numbers **a₁, a₂, a₃, a₄** at random between 0 and **p-1**
 - ▶ Hashing a key **k**
 - ▶ break **k** into **4** parts
 - ▶ **k₁, k₂, k₃, k₄**
 - ▶ output
 - ▶ output
- $$h(k) = \sum_{i=1}^4 a_i \cdot k_i \mod p$$
- ▶ Setup to store **150** students
 - ▶ choose **p=151**
 - ▶ choose **a₁=12, a₂=43, a₃=105, a₄=83**
 - ▶ Hashing a key **k=00238918**
 - ▶ break **k** into **k₁=00, k₂=23, k₃=89, k₄=18**
 - ▶ output
- $$h(00238918) = 50$$

HashTable with UHFs

- ▶ Hash table w/ chaining using a universal hash function family
 - ▶ Worst-case runtime of Get is $O(n)$ 
 - ▶ But UHFs guarantee that worst-case happens very *rarely*
 - ▶ We can “expect” that Get will have runtime $O(1)$
- ▶ What do we mean by expect?
 - ▶ remember that with UHFs we picked one function from family at random
 - ▶ in example we picked the values (a_1, a_2, a_3, a_4) at random
 - ▶ but for some functions in the family, keys will be well-spread & for others keys may be clustered
 - ▶ but if we were to compute the runtime of Hash Table with h a million times, where each time we sample a hash function at random from the family...
 - ▶ ...then the average of those runtimes would be $O(1)$
 - ▶ This is called “expected running time”

HashTable with UHFs

- ▶ Hash table w/ chaining using a universal hash function family
 - ▶ We can “expect” that Get will have runtime $O(1)$
- ▶ What do we mean by expect?
 - ▶ remember that with UHFs we picked one function from family at random
 - ▶ in the example we picked the values (a_1, a_2, a_3, a_4) at random
 - ▶ for some functions in the family, keys will be well-spread...
 - ▶ ...while for others keys will be poorly spread, e.g., all mapped to same value
 - ▶ but if we were to compute the runtime of Hash Table with a million times, where each time we sample a hash function at random from the family...
 - ▶ ...then the average of those runtimes would be $O(1)$
 - ▶ This is called “expected running time”

Why does Universal Hashing Work?

- ▶ See Chapter 1.5.2 in Dasgupta et al.
 - ▶ and/or read the proof in the following slides
 - ▶ You do not need to know the proof!

Proof of Universal Hashing

Inverses

- ▶ What is the inverse of a fraction x/y ?
 - ▶ y/x because $(x/y)(y/x)=1$
 - ▶ inverse is whatever we need to multiply it by to get 1
- ▶ What is the inverse of an int x (not 1)?
 - ▶ $1/x$ because $(x)(1/x)=1$
- ▶ What is the “integer” inverse of an int x (not 1)
 - ▶ there is none...
 - ▶ you can’t multiply an int w/ another int to get 1 (unless 1)

Modular Arithmetic

- ▶ If working modulo some number
 - ▶ Integers can have integer inverses!
- ▶ ex: let's work **mod 7**
 - ▶ inverse of 2 **mod 7** is 4 because $2 \times 4 \bmod 7 = 1$
 - ▶ inverse of 5 **mod 7** is 3 because $5 \times 3 \bmod 7 = 1$
- ▶ Is this always true?
 - ▶ ex: does 2 have an inverse **mod 4**?
 - ▶ $2 \times 0 \bmod 4 = 0; 2 \times 1 \bmod 4 = 2$
 - ▶ $2 \times 2 \bmod 4 = 0; 2 \times 3 \bmod 4 = 2$
 - ▶ No!
- ▶ But it is true when we work modulo a prime number
 - ▶ mod a prime, every number except 0 has a unique inverse

Analysis

- ▶ Prime **p** is the size of array
- ▶ x_1, x_2, x_3, x_4 are a banner ID in chunks
- ▶ y_1, y_2, y_3, y_4 are another banner ID in chunks
- ▶ If IDs are different, at least **1** of the chunks are diff
- ▶ Let's assume (wlog) it is the last one so
 - ▶ $x_4 \neq y_4$
- ▶ What is the probability that
 - ▶ $h(x_1, x_2, x_3, x_4) = h(y_1, y_2, y_3, y_4)$

Analysis

- ▶ What is the probability that
 - ▶ $h(x_1, x_2, x_3, x_4) = h(y_1, y_2, y_3, y_4)$
- ▶ Step #1:
 - ▶ find equivalent formulation of event
 - ▶ that makes the randomness explicit
 - ▶ what is the randomness here?
- ▶ Step #2:
 - ▶ what is probability of equivalent formulation?

Step I: Equivalent Formulation

$$h(x_1, x_2, x_3, x_4) = h(y_1, y_2, y_3, y_4)$$

by definition

$$a_1x_1 + \cdots + a_4x_4 \equiv a_1y_1 + \cdots + a_4y_4 \pmod{p}$$

move things

$$a_4x_4 - a_4y_4 \equiv (a_1y_1 + a_2y_2 + a_3y_3) - (a_1x_1 + a_2x_2 + a_3x_3) \pmod{p}$$

just some number;
let's call it c

$$a_4 \cdot (x_4 - y_4) \equiv c \pmod{p}$$

$$a_4 \equiv c \cdot (x_4 - y_4)^{-1} \pmod{p}$$

Step 2: Probability of Equiv. Formulation

- ▶ So hashes are equal when

$$a_4 \equiv c \cdot (x_4 - y_4)^{-1} \pmod{p}$$

- ▶ But

- ▶ x_4 and y_4 are different so $x_4 - y_4 \neq 0$
- ▶ and p is prime
- ▶ so $(x_4 - y_4)$ has unique inverse mod p
- ▶ So $c(x_4 - y_4)^{-1}$ can only take on one value
 - ▶ therefore a_4 can only take on one value
- ▶ What is the probability a_4 takes on that value?
 - ▶ a_4 is randomly chosen from p possible values so probability is $1/p$

Putting it all Together

- ▶ Prob. that some ID will collide w/ another ID
 - ▶ $1/p = 1/151$
- ▶ For some ID,
 - ▶ expected # of collisions w/ all other IDs is
 - ▶ $149/151 = 0.986\dots$
- ▶ Expected size of an ID's bucket is
 - ▶ $1 + 0.986\dots = 1.986\dots = O(1)$

End of Universal Hashing Proof

Summary

- ▶ Array-based Dictionaries
 - ▶ Add is worst-case $O(n)$
 - ▶ Get is worst-case $O(n)$
- ▶ Hash Table-based Dictionaries with UHFs
 - ▶ Add is
 - ▶ worst-case $O(n)$ but expected $O(1)$
 - ▶ Get is
 - ▶ worst-case $O(n)$ but expected $O(1)$

Q: what can we build from dictionaries?

A (Basic) Search Engine

- ▶ Build a dictionary that maps keywords to URLs
 - ▶ query dictionary on keyword to retrieve URLs
- ▶ In context of search engines
 - ▶ the dictionary is often called an *Index*

A (Basic) Search Engine

- ▶ For each keyword **word** w/ a list of relevant URLs $\text{url}_1, \dots, \text{url}_m$
 - ▶ store the pairs $(\text{word} | 1, \text{url}_1), \dots, (\text{word} | m, \text{url}_m)$ in a dict **Index**
 - ▶ where “|” is string concatenation
 - ▶ Store the pair (word, m) in an auxiliary dictionary **Counts**
- ▶ To search for a keyword **Brown**
 - ▶ retrieve the count for **Brown** by querying **Count.get(Brown)**
 - ▶ to recover URLs, query **Index** on keys $\text{Brown} | 1, \dots, \text{Brown} | m$
 - ▶ **Index.get(word|1), ..., Index.get(word|m)**

Build Index

```
function build_index(page_list):
    index = dict()
    counts = dict()
    for page in page_list:
        for word in page:
            try:
                count = counts.get(word)
            except KeyError:
                counts.put(word, 0)
                count = counts.get(word)
            counts.put(word, counts[word] + 1)
            key = word + str(counts.get(word))
            index.put(key, page.url)
    return index
```

- ▶ `build_index` is $O(nm)$ time
 - ▶ where **n** is number of pages and **m** is maximum number of words per page

Search Index

```
def search_index(index, word):
    output_list = list()
    count = 1
    while True:
        try:
            url = index.get(word + str(count))
            count = count + 1
        except KeyError:
            break
        output_list.append(url)
    return output_list
```

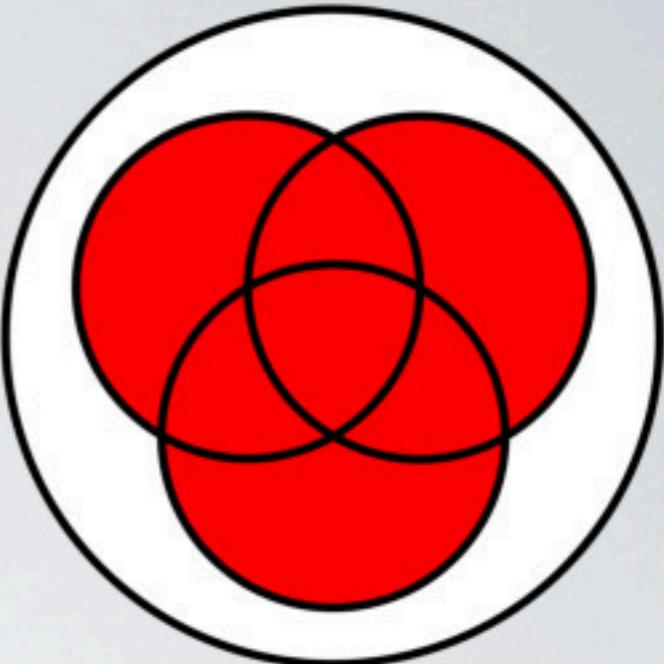
- ▶ If dictionary is implemented with hash table
 - ▶ `search_index` is expected $O(1)$ time
 - ▶ fast no matter how many pages and words

A (Basic) Search Engine

- ▶ What's missing from our “search engine”?
 - ▶ No ranking!
 - ▶ But we'll learn how to rank later in the course
 - ▶ ...after we learn about graphs

Sets

- ▶ Collection of elements that are
 - ▶ distinct and unordered
 - ▶ ...unlike lists and arrays



Set ADT

- ▶ **add(object):**
 - ▶ adds object to set if not there
- ▶ **remove(object):**
 - ▶ removes object from set if there
- ▶ **boolean contains(object):**
 - ▶ checks if object is in set
- ▶ **int size():**
 - ▶ returns number objects in set
- ▶ **boolean isEmpty():**
 - ▶ returns TRUE if set is empty; FALSE otherwise
- ▶ **list enumerate():**
 - ▶ returns list of objects in set (in arbitrary order)



Set Data Structure

- ▶ How can we implement a Set?
- ▶ Using an *expandable array*
- ▶ add: $O(1)$
- ▶ contains: $O(n)$ (scan array)
- ▶ remove: $O(n)$ (find & compress)
- ▶ Can we do better?



Sets from Hash Tables

- ▶ We can implement sets with a hash table
- ▶ Sometimes called a Hash Set

```
function add(object):  
    index = h(object)  
    table[index].append(object)
```

Expected $O(1)$

```
function contains(object):  
    index = h(object)  
    for elt in table[index]:  
        if elt == object:  
            return true  
    return false
```

Expected $O(1)$

HashMap vs. HashSet

- ▶ HashMap
 - ▶ Hash table implementation of a dictionary
- ▶ HashSet
 - ▶ Hash table implementation of a set